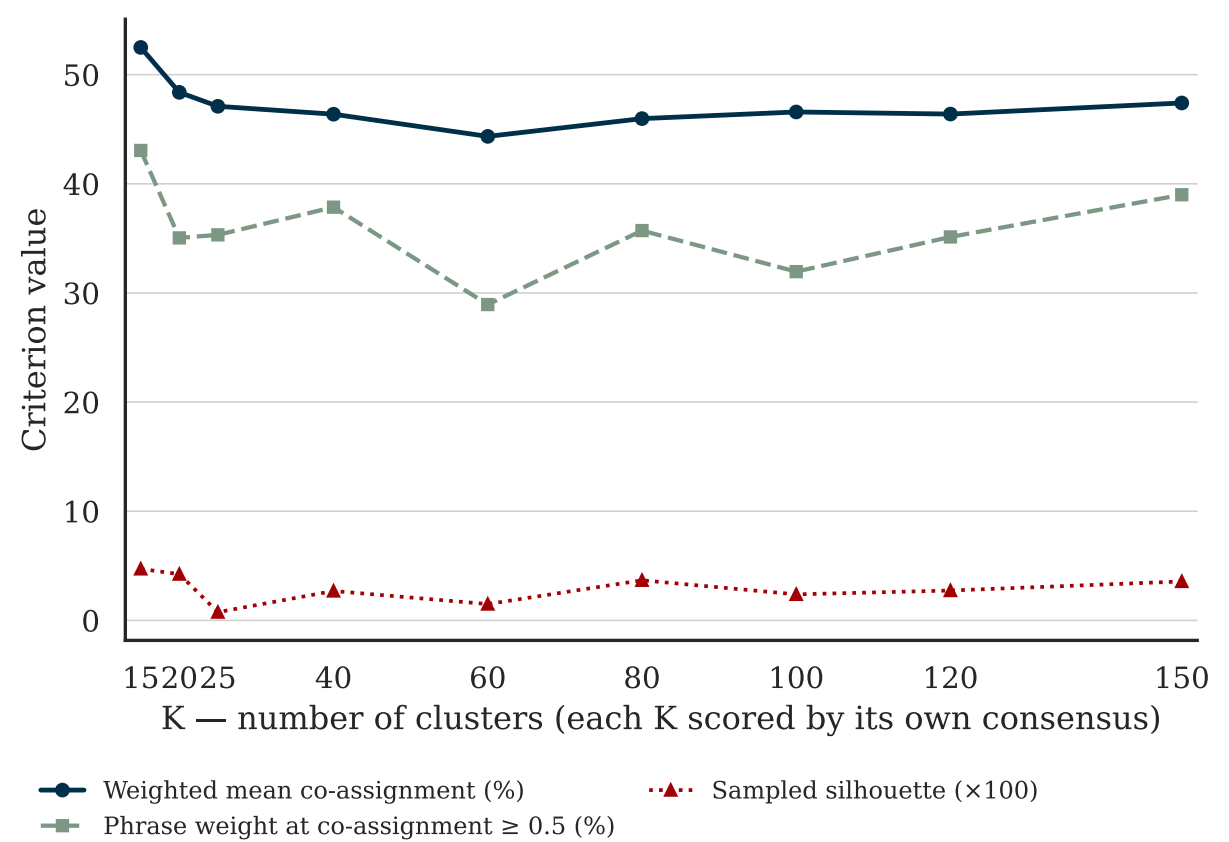

Note: consensus clustering of the Track-1 (hearing-level) topic phrases, weights = mention counts. For each candidate K, 30 bootstrap replicates (weight-proportional 80% subsamples) are clustered at that K; average linkage on the per-K co-assignment matrix, cut at K, gives the clusters scored here. Stability is reported per cluster in the codebook and gates nothing — the 0.5-bar series is a reference only.