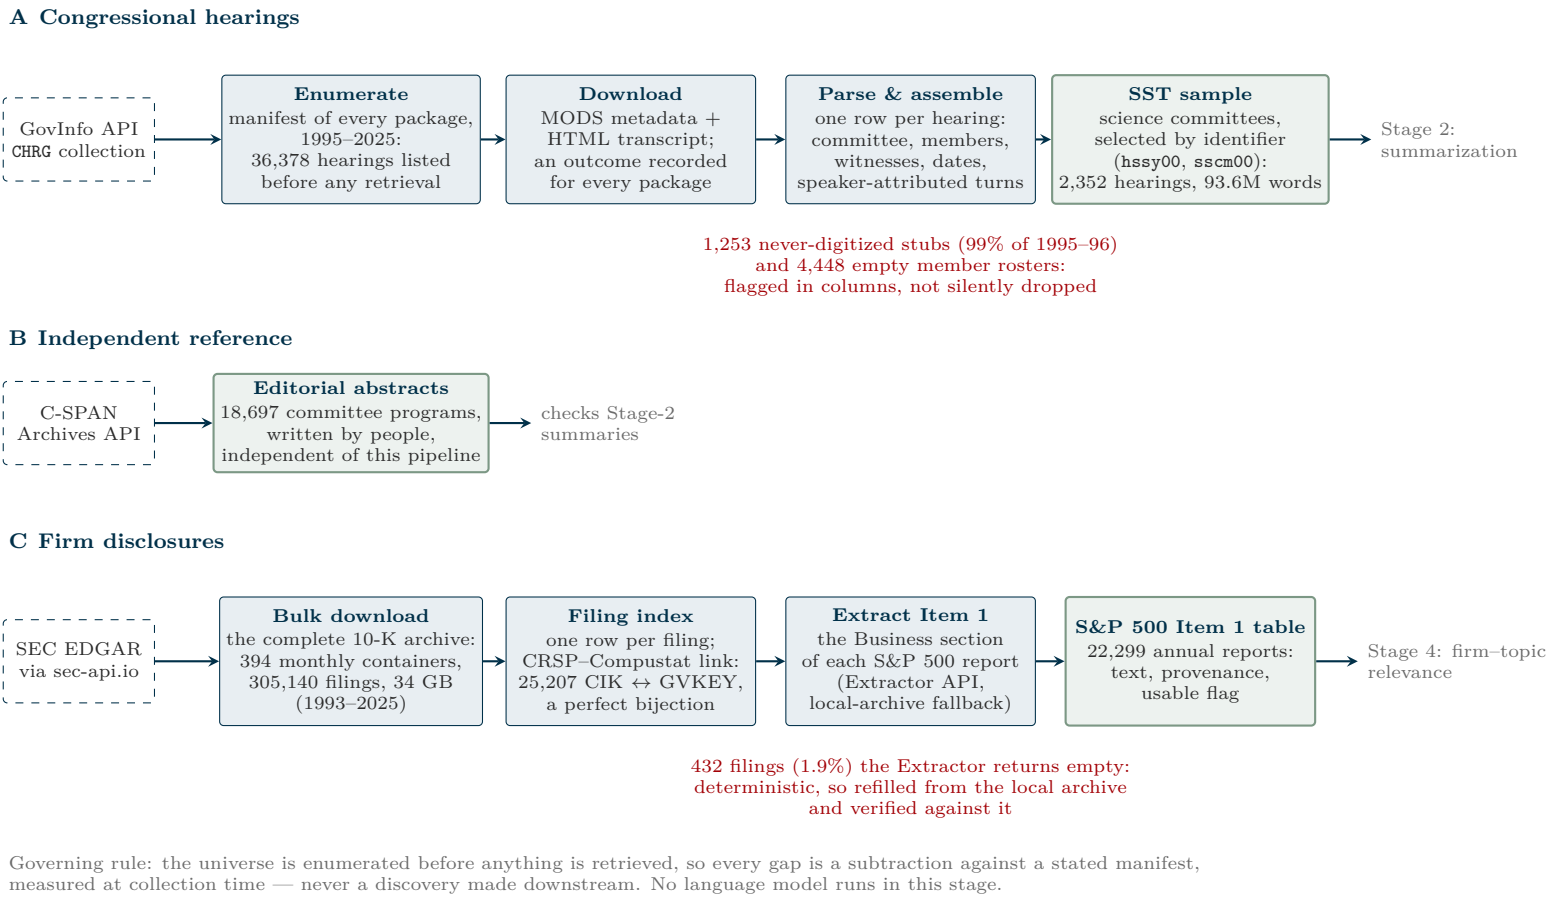
A Congressional hearings
GovInfo API
CHRG collection
Enumerate
manifest of every package, 1995–2025: 36,378 hearings listed before any retrieval
Download
MODS metadata + HTML transcript; an outcome recorded for every package
Parse & assemble
one row per hearing: committee, members, witnesses, dates, speaker-attributed turns
SST sample
science committees, selected by identifier (hssy00, sscm00): 2,352 hearings, 93.6M words
Stage 2: summarization
1,253 never-digitized stubs (99% of 1995–96) and 4,448 empty member rosters: flagged in columns, not silently dropped
B Independent reference
C-SPAN Archives API
Editorial abstracts
18,697 committee programs, written by people, independent of this pipeline
checks Stage-2 summaries
C Firm disclosures
SEC EDGAR via sec-api.io
Bulk download
the complete 10-K archive: 394 monthly containers, 305,140 filings, 34 GB (1993–2025)
Filing index
one row per filing; CRSP–Compustat link: 25,207 CIK ↔ GVKEY, a perfect bijection
Extract Item 1
the Business section of each S&P 500 report (Extractor API, local-archive fallback)
S&P 500 Item 1 table
22,299 annual reports: text, provenance, usable flag
Stage 4: firm–topic relevance
432 filings (1.9%) the Extractor returns empty: deterministic, so refilled from the local archive and verified against it
Governing rule: the universe is enumerated before anything is retrieved, so every gap is a subtraction against a stated manifest, measured at collection time — never a discovery made downstream. No language model runs in this stage.